

Peter Cotton

Chief Data Scientist
Intech Investments

Microprediction

Collective Intelligence in Real-time

August 5, 2020



Outline

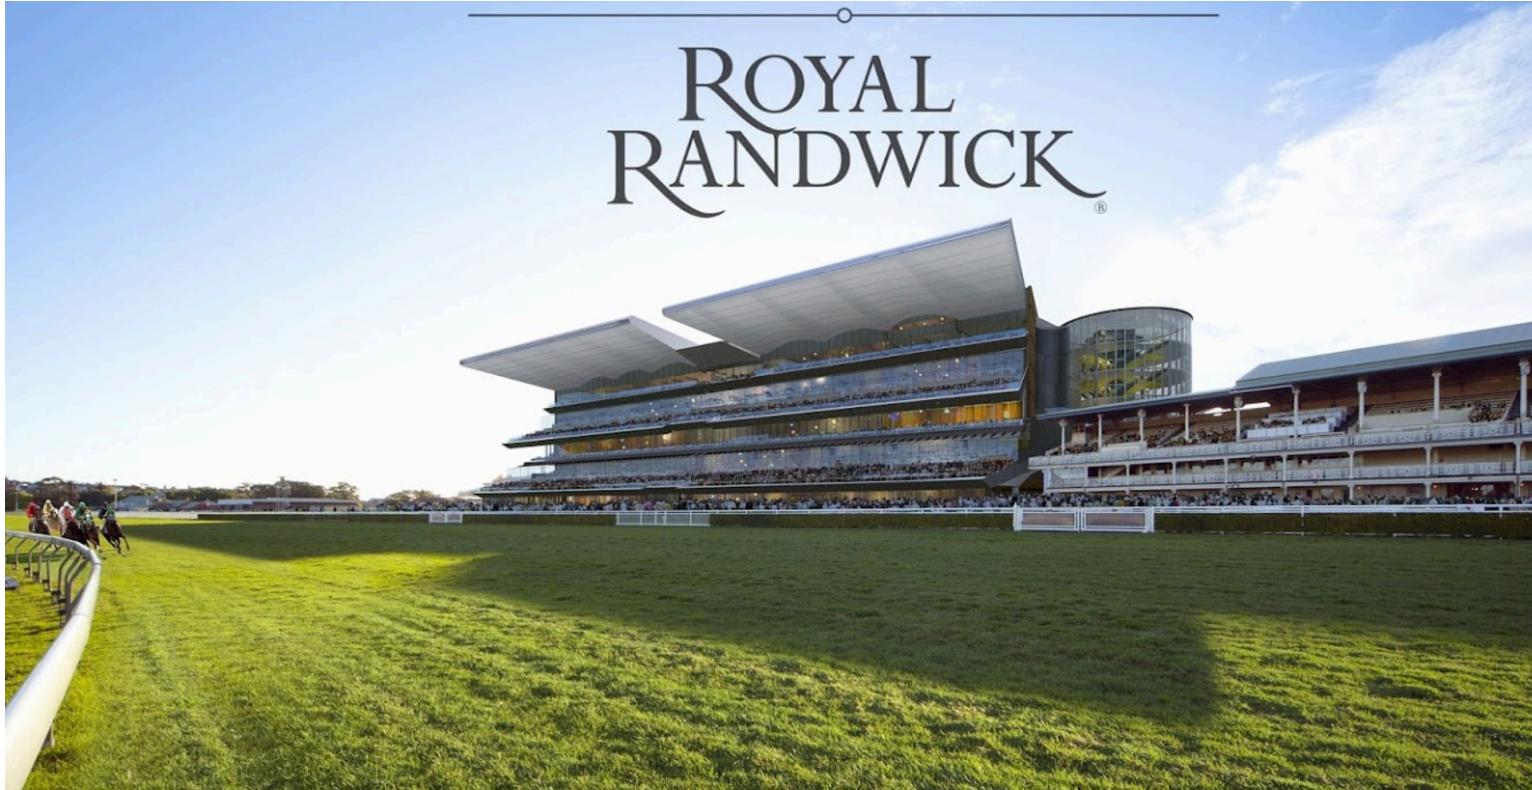
- Motivation for Microprediction.org
 - A little about how my experiences led me to try to create a prediction web
- How you can use it
 - Creating a stream
 - Predicting a stream
 - How to create and modify crawlers

Mission

- Turnkey bespoke repeated prediction for everyone

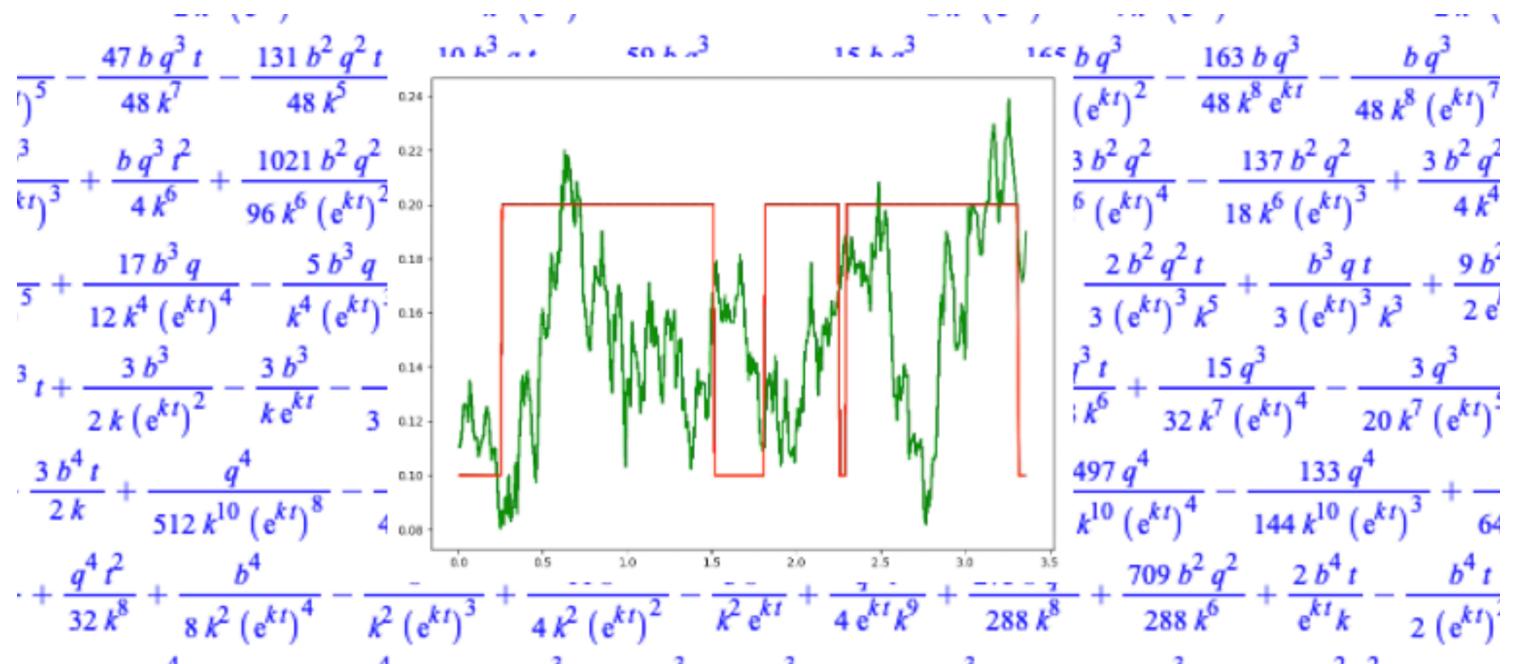
My background: gambler

In college, worked for a professional gambler who made his living on pari-mutuels



My background: student

Doctorate on PDEs and homogenization. (George Papanicolaou)



<https://www.linkedin.com/pulse/got-milk-homogenization-survival-probability-peter-cotton-phd/>

My background: quant

First job at Morgan Stanley 2001-2007.
Invented closed form pricing of Collateralized Debt Obligations (CDOs)

WIRED

Recipe for Disaster: The Formula That Killed Wall Street

FELIX SALMON 02.23.09 12:00 PM

Share

 SHARE

 TWEET

 COMMENT

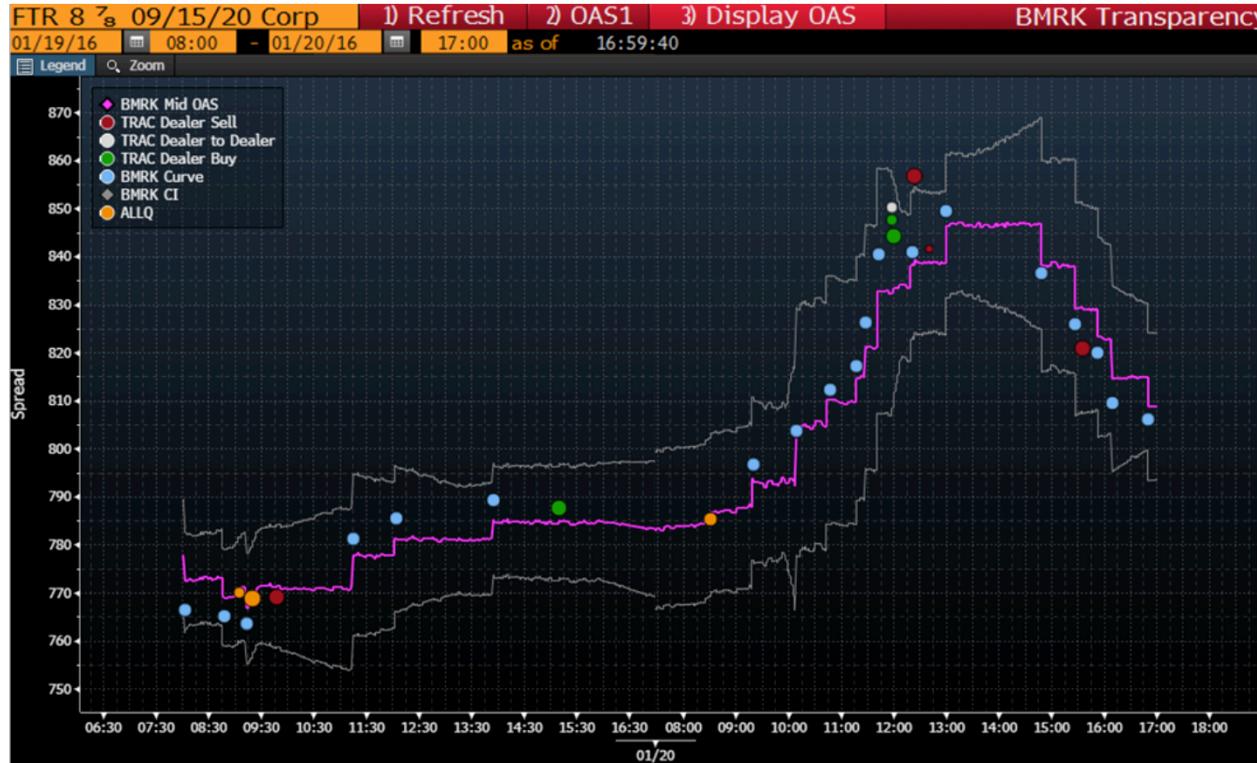
 EMAIL

Recipe for Disaster: The Formula That Killed Wall Street



My background: entrepreneur

Founded a company in 2009 that pioneered large scale data assimilation and Bayesian filtering of the U.S. corporate credit market (eaten by BBG 2013)



The BMRK Transparency screen displays real-time movements in credit spreads based on market data (trades and quotes) and synthetic observations from BMRK's issuer curve model.

My background: quant again

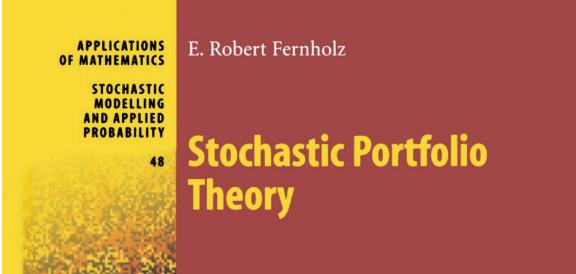
Worked at J.P. Morgan. Introduced the use of optimal control theory for over-the-counter trading of corporate bonds.

$$\begin{aligned}\frac{\tau c(x)}{s} &= p \sup_{m^\uparrow} \left\{ \overbrace{(m^\uparrow - K^\uparrow(x; s))}^{\text{net gain}} P^\uparrow(m^\uparrow) \right\} \\ &\quad + (1 - p) \sup_{m^\downarrow} \{ (m^\downarrow - K^\downarrow(x; s)) P^\downarrow(m^\downarrow) \} \\ &\quad - p \sup_{m^\uparrow} \{ (m^\uparrow - K^\uparrow(0; s)) P^\uparrow(m^\uparrow) \} \\ &\quad - (1 - p) \sup_{m^\downarrow} \{ (m^\downarrow - K^\downarrow(0; s)) P^\downarrow(m^\downarrow) \}\end{aligned}$$

Also initiated J.P. Morgan's privacy preserving computation and crowdsourcing initiatives.

My current role

Intech Investments



RANKED AS A LEADING QUANT MANAGER BY AUM¹

4 / GLOBAL EQUITY # 5 / U.S. EQUITY # 5 / ENHANCED EQUITY # 6 / DEFENSIVE EQUITY



<http://finmathblog.blogspot.com/2012/12/excess-return-on-portfolio-of-lognormal.html>



Microprediction – What is it?

Constantly repeated prediction



Artificial Intelligence – What is it?

Constantly repeated prediction

(of value functions)

Microprediction – What is it not?

Predicting next year's GDP



Microprediction (including Crislano's Questions)

- How to improve demand forecasting? (yes, maybe but ...hmmm...)
- How to improve short demand forecasting? (yes ... better!)
- How to predict call center volumes? (yes)
- How to predict software failure? (yes, maybe, but ...)
- How to forecast stock price? (yes, everything but the mean!)
- How to predict value functions (RL)? (yes, why not?)
- How to predict electricity prices? (yes)
- How to predict bakery sales? (yes)
- How to predict activity near New York city hospitals? (yes)
- How to predict the winner of an election? (yes, but indirectly)

Someone said to me once...

Hey this thing is cool but if you could provide an API that would price any bond you give it, I will fund you!

Subsequently I got to pondering...

What about an API that predicts anything?

So here it is

Community Microprediction API 0.1

[Base URL: /]

<https://devapi.microprediction.org/swagger.json>

Hello. Did you know there is a [human readable site](#) and numerous [articles](#) about these APIs?

live Publish a live source of data, or retrieve data.

GET `/live/{name}` Retrieve value or derived value

PUT `/live/{name}` Publish scalar value

PATCH `/live/{name}` Modify time to live

DELETE `/live/{name}` Delete a stream



Example: Predicting water levels for NOAA

```
from credentials import ELFEST_BOBCAT # You'll have to supply your own
import pandas as pd
from microprediction import MicroWriter
import time

# Video explanation of this example:
# https://vimeo.com/443203883

mw = MicroWriter(write_key=ELFEST_BOBCAT) # See creating_a_key.py
STREAM_NAME = 'water.json'

def water_height():
    df = pd.read_csv('https://www.ndbc.noaa.gov/data/realtimedata/21413.dart')
    return float(df.iloc[1, :].values[0].split(' ')[-1])

def poll_for_an_hour():
    for _ in range(4):
        mw.set(name=STREAM_NAME, value=water_height())
        time.sleep(15 * 60)
        print('.', flush=True)

if __name__ == '__main__':
    poll_for_an_hour()
```

<https://www.linkedin.com/feed/update/urn:li:activity:6694740397144596480/>

How does it work?

You make it work!

```
from microprediction import MicroCrawler

if __name__ == '__main__':
    crawler = MicroCrawler(difficulty=9)
    crawler.run()
```

Run this right now



How does a crawler do?

70, 310, 910, 3555



Sends 225 guesses of the next data point to arrive after ___ seconds have expired

Maintains a TODO list of what to predict, when

Performance

```
+35.7987: 70::z1~die~70
+28.3926: 3555::z1~copula_z~70
+24.7674: 70::z1~size~3555
+24.1433: 3555::z1~btc_eur~70
+22.064: 70::z1~bart_delays~70
+16.8854: 3555::z1~die~3555
+15.1565: 70::z1~coin_a~3555
+15.1038: 3555::fcx
+13.6464: 910::z1~copula_y~70
+13.0955: 3555::z1~copula_z~3555
+10.9293: 310::z1~bart_delays~3555
+10.4233: 3555::z1~ma~70
+9.7736: 310::z1~traffic_speed_deltas~70
+9.6291: 910::z1~bart_delays~3555
+9.0671: 910::z1~coin_b~70
+8.9786: 3555::z1~die~70
+8.8225: 70::z1~coin_b~70
+8.0184: 3555::z1~copula_y~3555
+7.3876: 310::z1~fcx~3555
+5.3377: 910::z1~fcx~70
+3.8892: 910::z1~helicopter_psi~70
+3.8807: 310::z1~dvn~70
+3.2549:
910::z1~South_Australia_Electricity_Price~3555
+3.056: 310::z1~traffic_absolute_speed~70
+2.9035: 910::z1~cat~3555
+2.6255: 310::z1~goog~70
+2.1289: 70::z1~seattle_wind_speed~70
+2.1007: 910::z1~c5_bitcoin~3555
+2.0632: 3555::z1~dvn~3555
+2: 910::z1~South_Australia_Electricity_Price~70
+0.4173: 910::z1~seattle_wind_direction~3555
+0: 310::z1~copula_x~70
+0: 70::z1~copula_z~70
+0: 70::z1~helicopter_theta~3555
+0: 70::coin_a
+0: 910::z1~copula_x~3555
-0.4545: 3555::z1~traffic_speed_deltas~3555
-1.0507: 310::z1~seattle_wind_direction~70
-1.051: 910::z1~copula_z~3555
-1.3249: 3555::c5_bitcoin
-1.5616: 3555::usmv
-2.1537: 70::seattle_wind_direction
-2.7451: 910::z1~c5_ethereum~70
-3.081: 310::z1~sq~70
-3.4996: 310::seattle_wind_direction
-3.6291: 70::z1~helicopter_psi~3555
-4.455: 910::z1~c5_ripple~3555
-4.8213: 910::z1~helicopter_theta~70
-5.2059: 3555::cat
-5.6715: 910::z1~c5_cardano~3555
-6.136: 910::z1~quasi~70
```

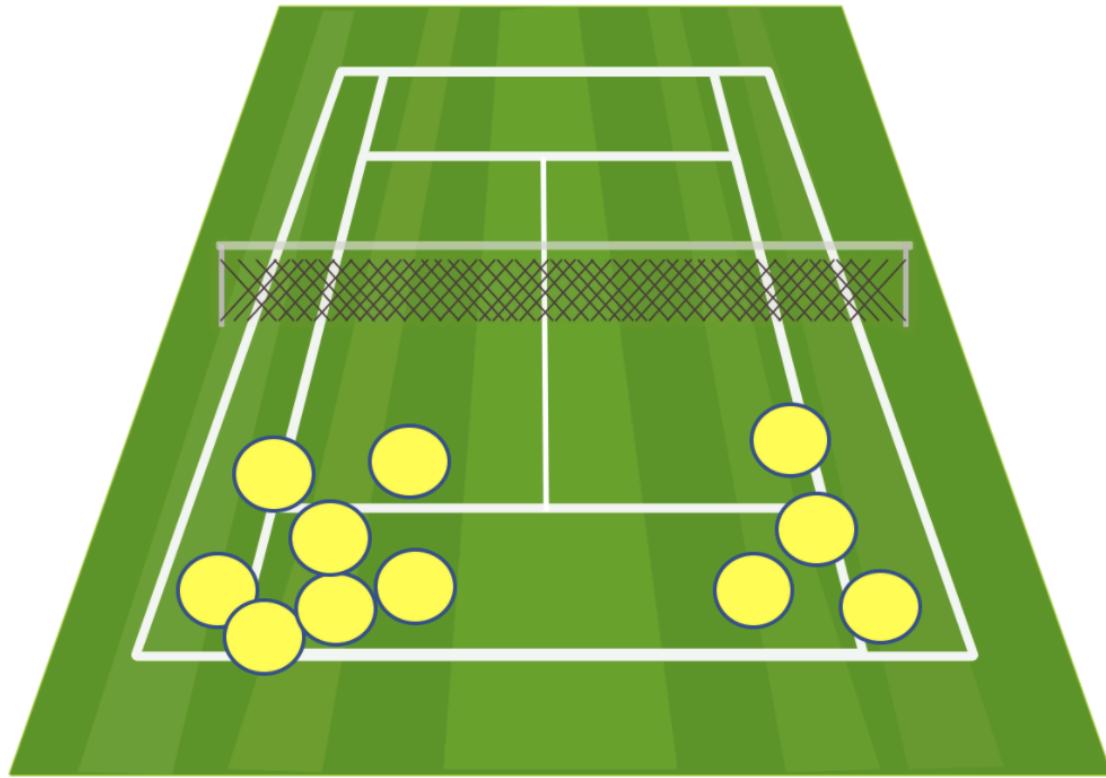
Searches for new data streams

Gives up if performance is poor

-0.33333 from z1~three_body_x~70
Awarded for 2 / 1125 submissions
Time: 2020-08-05 03:16
Delay: 3555
Stream Owner: cellosebobcat...

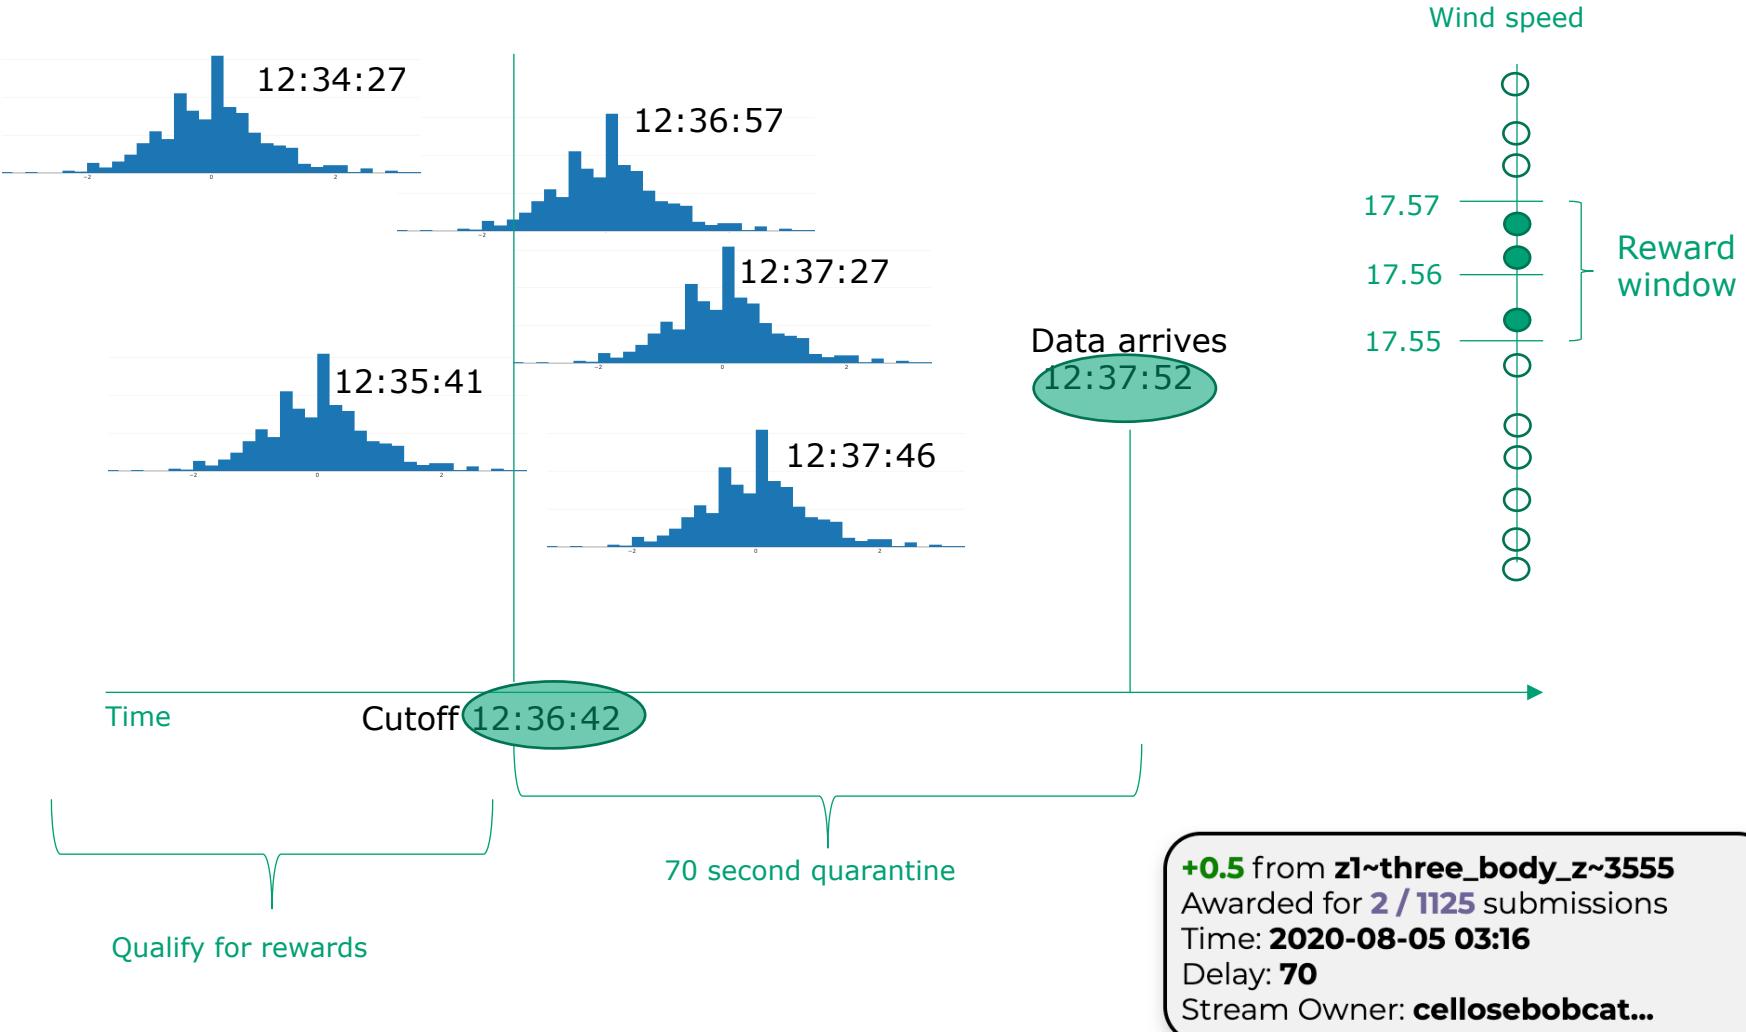


Why send 225 points instead of just the mean?



Rewarding of quarantined distributional predictions

System timestamps incoming distributional predictions
Rewards guesses in a neighborhood of the true value



Crawlers visit lots of streams

Microprediction Publish Data ▾ Submit Predictions ▾ Data Streams ▾ Learn More ▾

Dashboard

View all information associated with your write key, including your active streams, performance, and transactions. **Information loaded from `/home/{write_key}`**

Your write key Submit Account

Overview
e86a1ableca7441c1034b5adb4c22844
Memorable ID: Exhalable Cat
Balance: -685.7878
Distance to Bankruptcy: 3410.21

Active Streams
3555::z1~btc_eur~70
70::copula_z
910::copula_z
3555::copula_z
910::usmv
3555::usmv
310::badminton_x
70::z1~helicopter_theta~3555
70::seattle_wind_direction

Performance
+27.0033: 70::z1~die~70
+23.1176: 70::z1~bart_delays~70
+14.3214: 3555::z1~die~3555
+13.4026: 3555::z1~die~70
+13.2578: 310::z1~pandemic_recovered~70
+11.3291: 910::z1~pandemic_infected~3555
+11.2079: 310::z1~fcx~3555
+10.6883: 70::z1~coin_a~3555
+10.4233: 3555::z1~ma~70
+10.3136: 910::z1~bart_delays~3555
+9.7732: 910::z1~helicopter_theta~70
+8.823: 910::z1~copula_y~70
+8.125: 3555::z1~copula_z~70
+6.8122: 310::z1~btc_aud~70
+6.7991: 3555::fcx
+6.5304: 910::z1~coin_b~70
+6.4384: 310::z1~size~70
+6.0024: 3555::z1~btc_eur~70

Transactions
-1: z1~traffic_absolute_s
-0.5: z1~badminton_y~3
+0.375: z1~badminton_x
+1: z1~badminton_y~70
-0.5: badminton_x
-0.3333: z1~copula_z~3555
+0.6667: z1~copula_z~3
+0.1111: z1~copula_y~3555
+0: z1~copula_x~3555
+2: z1~copula_x~3555

Errors
No Errors

It is easy to incorporate any time series model

A model is just something that spits out 225 guesses of a future value

```
class BenchmarkCrawlerExample(SimpleCrawler):

    def __init__(self, **kwargs):
        super().__init__(**kwargs)

    def sample(self, lagged_values, lagged_times, **ignored):
        model.fit(lagged_values)
        half_width = 0.5 / self.num_predictions
        prctls = np.linspace(half_width, 1 - half_width, self.num_predictions)
        return [model.invcdf(p) for p in prctls]
```

There are other abstractions (classes) you can use

Example: A t-digest used as a `DistributionMachine` inside `SequentialStreamCrawler`

```
from microprediction.config_private import STATESBOY_CAT
from tdigest import TDigest
from microprediction import SequentialStreamCrawler, DistributionMachine


class DigestMachine(DistributionMachine):

    def __init__(self):
        super().__init__()
        self.digest = TDigest()

    def update(self, value, dt=None, **ignored):
        self.digest.update(value)

    def inv_cdf(self, p):
        return self.digest.percentile(100. * p)

if __name__ == "__main__":
    crawler = SequentialStreamCrawler(write_key=STATESBOY_CAT, min_lags=500, machine_type=DigestMachine)
    crawler.set_repository(
        url='https://github.com/microprediction/microprediction/blob/master/crawler_examples/statesboy_cat.py')
    crawler.min_lags = 500
    crawler.run()
```

For more examples of making crawlers...

The screenshot shows a GitHub repository page for the user `microprediction`. The repository name is `microprediction / microprediction`. The `Code` tab is selected. There are 6 issues and 6 pull requests. The `Actions`, `Projects`, `Wiki`, and `Security` tabs are also visible.

The current branch is `master`. The repository path is `microprediction / crawler_examples /`.

A pull request by `Peter Cotton` and `Peter Cotton` is shown, with the commit message: "More aggressive withdrawal if there is activity excess".

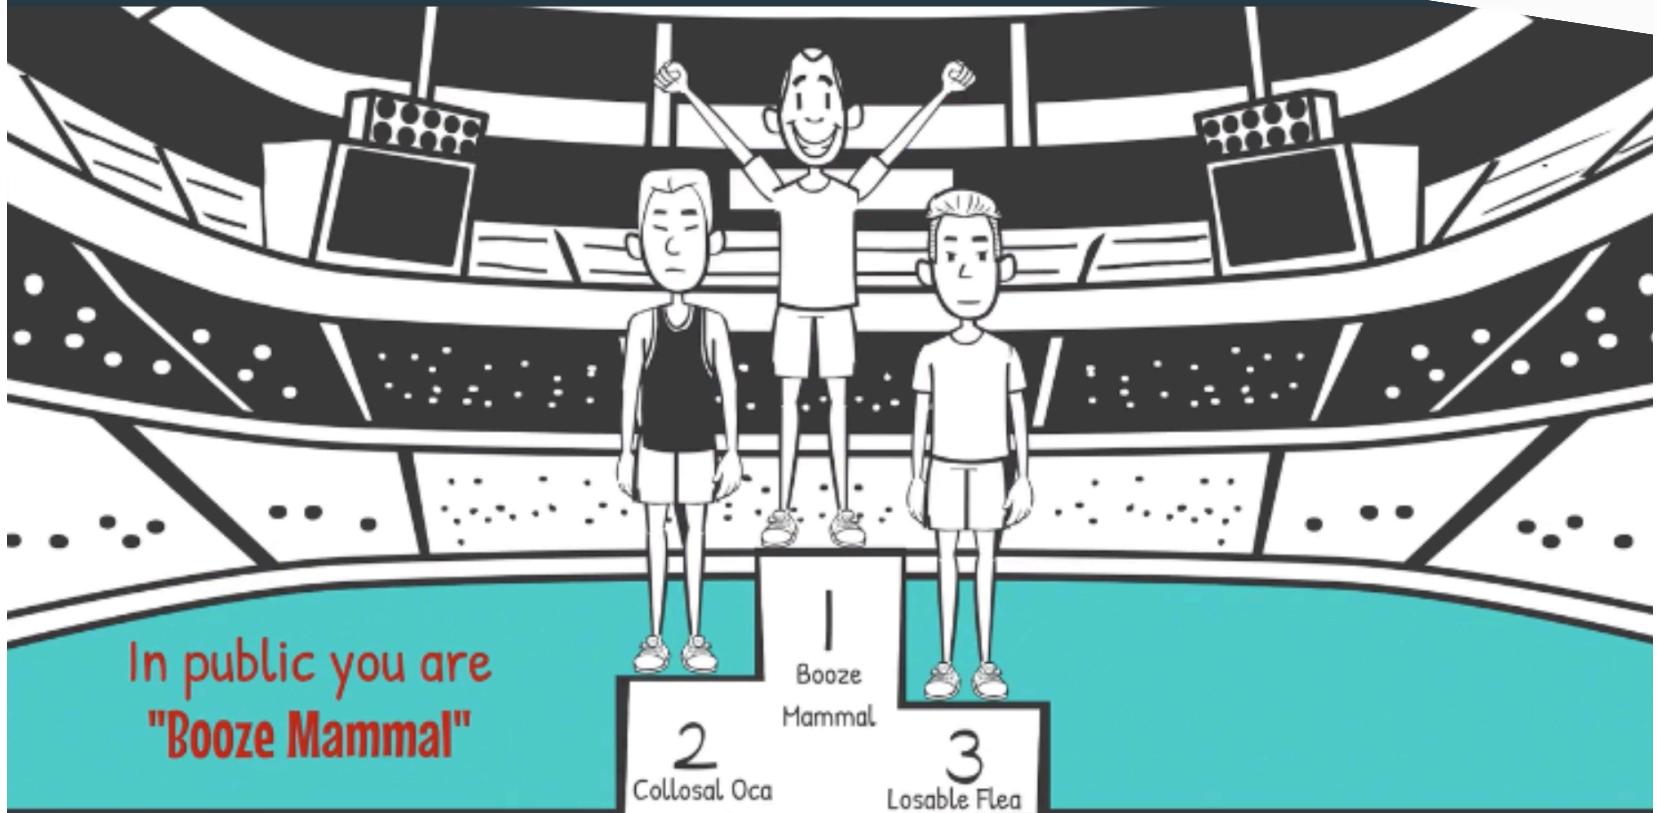
The list of files and their descriptions is as follows:

- `__init__.py`: stop loss and FlammableCod example
- `booze_mammal.py`: Added better timing
- `dale_leech.py`: Updated README with links to examples
- `decastyle_cat.py`: Added statefulcrawler
- `dirac_crawler.py`: Updated README with links to examples
- `flammable_cod.py`: Updated README with links to examples
- `simple.py`: Fixed minor bug in simplecrawler example
- `statesboy_cat.py`: More aggressive withdrawal if there is activity excess
- `thalloid_cat.py`: Added SequentialStreamCrawler

What's with the silly names?

No registration required!

```
hash: "76ecabee3c0ebd682e935f99a7c07dd4",
key: "f1823fc339f1a8f6d65ccd666e20ea97",
length: 8,
pretty: "Theca Bee"
```



Explained: <https://vimeo.com/397352413>

Click here to create a short MUID: <http://www.muid.org/create/>

Take another look ...

First thing the “hello world” crawler does is create a MUID

```
from microprediction import MicroCrawler

if __name__ == '__main__':
    crawler = MicroCrawler(difficulty=9)
    crawler.run()
```

The hash of the private identifier is memorable

```
hash: "76ecabee3c0ebd682e935f99a7c07dd4",
key: "f1823fc339f1a8f6d65ccd666e20ea97",
length: 8,
pretty: "Theca Bee"
```

... or at least some of it is after you convert 8->x, 7->t et cetera

What can crawlers do?

Anything you want them to

1. Find exogenous data
2. Use automated model search
3. Use the prediction API recursively
4. Optimize their navigation
5. Predict when data will arrive
6. Exploit online statistical methods
7. Specialize to certain types of time series
8. Specialize to implied time series (zscores, zcurves)
9. Reset their performance
10. Withdraw from streams where they lose balance

Rank	MUID	Points
1	Decastyle Cat	+9675.118
2	Flammable Cod	+3753.247
3	Semese Bat	+3359.224
4	Chatty Fly	+2040.697
5	Mesole Mammal	+1822.514
6	Carryover	+1532.995
7	Doddle Mammal	+1178.333
8	Behest Mammal	+979.33
9	Choosy Beetle	+777.924
10	Staboy Bat	+464.864



What can you do with a microprediction API?

Standardize data relative to competitive community forecasts

z1~fcx~70

[Go to Stream →](#)

Horizon: 70 sec 310 sec 910 sec 3555 sec

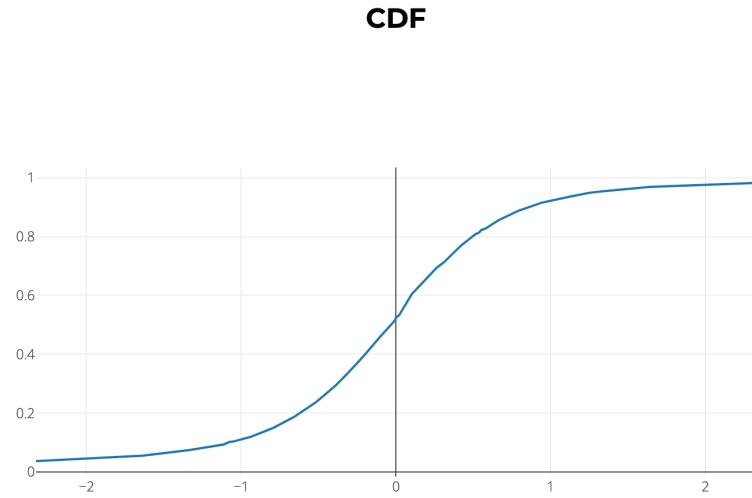
Leaderboard

Rank	MUID	Points
1	Comal Cheetah	+7.101
2	Decastyle Cat	+5.595
3	Exhalable Cat	+5.338
4	null	+0.8
5	Flammable Cod	-3.793
6	Cellose Bobcat	-15.04

Lagged Values

Timestamp	Z-Score
8/4 16:01:06	-0.51996
8/4 16:00:11	0.55508
8/4 15:59:18	-0.00448
8/4 15:58:05	-0.04721
8/4 15:57:04	1.48253
8/4 15:56:03	0.55323
8/4 15:55:05	-0.00446
8/4 15:54:03	-0.00445
8/4 15:53:04	0.58113
8/4 15:52:11	-1.29729
8/4 15:51:05	1.24723
8/4 15:50:06	-0.5483
8/4 15:49:04	-0.55082
8/4 15:48:04	-0.53286
8/4 15:47:04	-1.07452
8/4 15:46:06	-0.39177
8/4 15:45:13	0.26186

CDF



What can you do with a microprediction API?

Predict absolute values

traffic_absolute_speed

Live Current Value: **49.7**

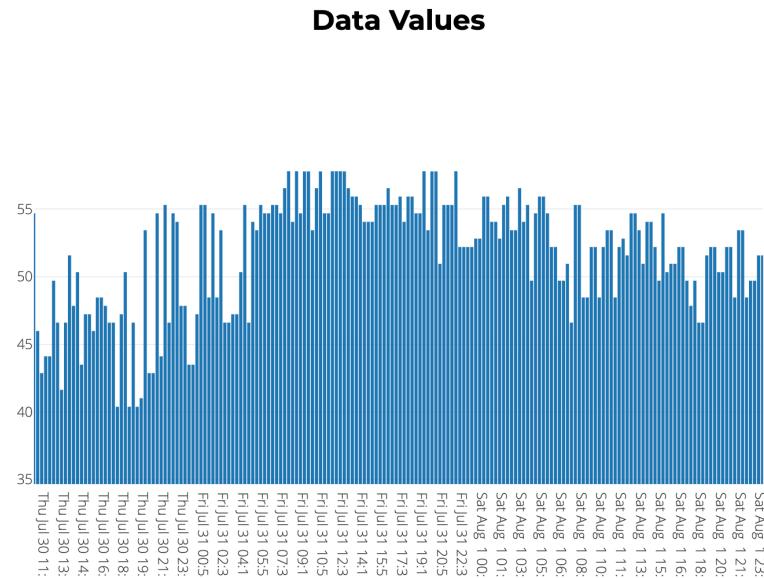
[← Go to Competitions](#) [← Go to Z1](#)

Leaderboard

Rank	MUID	Points
1	Staboy Bat	+231.03
2	Doodle Bee	+161.968
3	Chatty Fly	+101.565
4	Behest Mammal	+97.258
5	Decastyle Cat	-8.737
6	Lettable Clam	-8.816
7	Bedabble Toad	-9
8	Statesboy Cat	-10.129
9	Flammable Cod	-10.5
10		-10.000

Lagged Values

Timestamp	Data
8/4 22:39:32	49.7
8/4 22:19:33	49.08
8/4 21:59:33	49.7
8/4 21:39:32	49.08
8/4 21:19:32	49.7
8/4 20:59:34	49.08
8/4 20:39:32	49.08
8/4 20:13:40	54.05
8/4 19:53:40	36.03
8/4 19:33:41	50.95
8/4 19:13:39	49.7
8/4 18:53:42	45.98
8/4 18:34:07	49.7
8/4 18:13:44	45.98
8/4 17:53:41	45.98
8/4 17:33:48	47.22
8/4 17:13:41	50.33



What can you do with a microprediction API?

Predict changes in values

cop

Live Current Value: **1.0627**

[← Go to Competitions](#) [← Go to Z1](#)

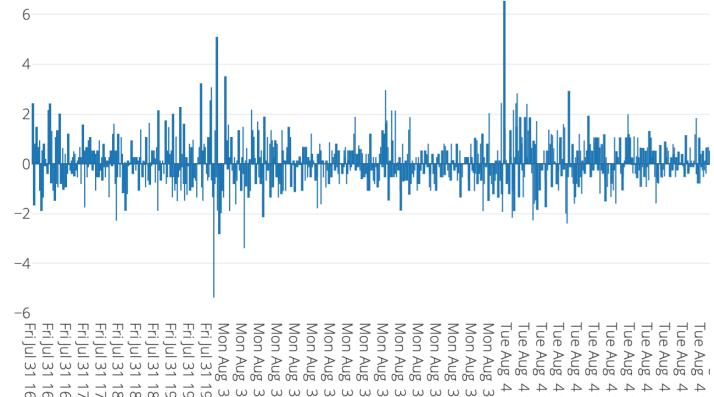
Leaderboard

Rank	MUID	Points
1	Cellose Bobcat	+44.149
2	Decastyle Cat	+31.662
3	Mesole Mammal	+11.565
4	null	+2.4
5	Bedabble Toad	-0.117
6	Bal Llama	-0.8
7	Lettable Clam	-3.291
8	Oft Sloth	-10.802
9	Exhalable Cat	-12.766

Lagged Values

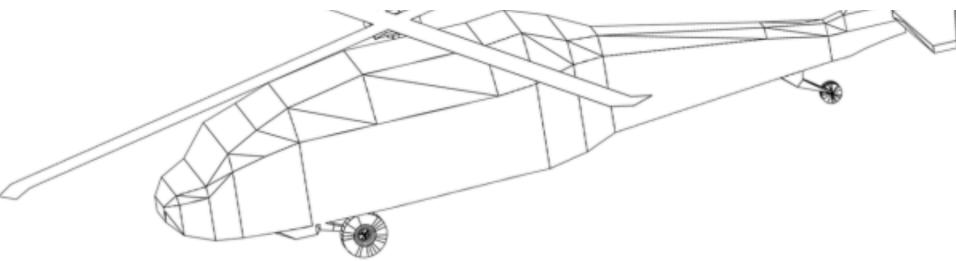
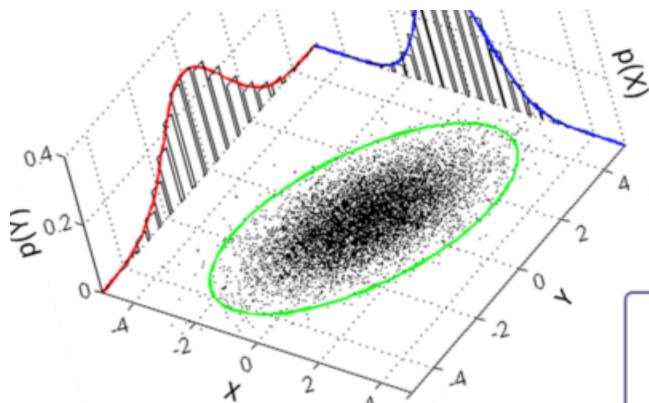
Timestamp	Data
8/4 16:01:05	1.0627
8/4 16:00:11	0.26585
8/4 15:59:18	0.53191
8/4 15:58:04	0.26606
8/4 15:57:03	0.13306
8/4 15:56:03	0.63094
8/4 15:55:04	-1.29605
8/4 15:54:03	1.86344
8/4 15:53:04	-0.66591
8/4 15:52:11	-1.19753
8/4 15:51:05	2.12993
8/4 15:50:06	-0.26649
8/4 15:49:04	-0.53277
8/4 15:48:04	-0.13315
8/4 15:47:04	-0.66547
8/4 15:46:06	-0.93091

Data Values



What can you do with a microprediction API?

Crowdsource copulas (bivariate)



helicopter_theta

Live Current Value: **-0.67467**

[← Go to Competitions](#) [← Go to Z1](#)

Leaderboard

Rank	MUID	Points
1
2
3
4
5

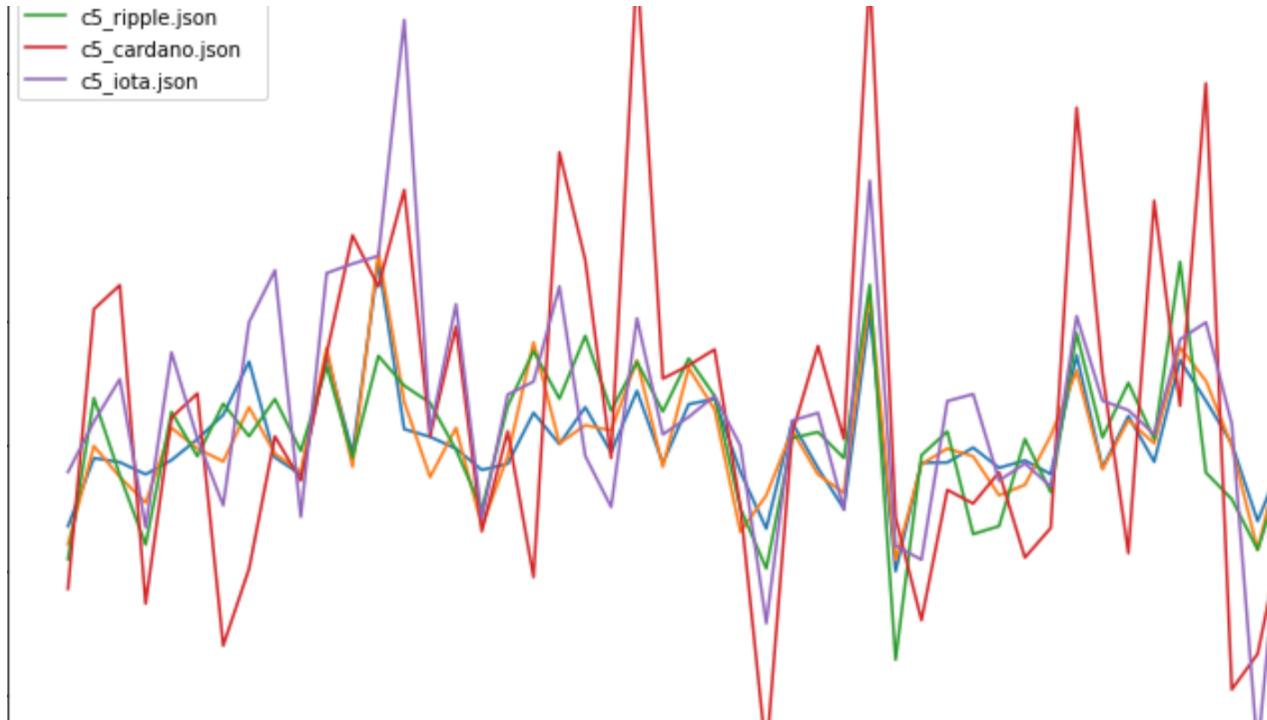
Lagged Values

Timestamp	Data
7/8 08:54:09	-0.67467
7/8 08:47:09	-0.67589
7/8 08:40:09	-0.67386
7/8 08:33:08	-0.68115
7/8 08:26:08	-0.68034
7/8 08:22:12	-0.68155

Explanation: <https://www.linkedin.com/pulse/helicopulas-peter-cotton-phd/>

What can you do with a microprediction API?

Crowdsource copulas (trivariate)



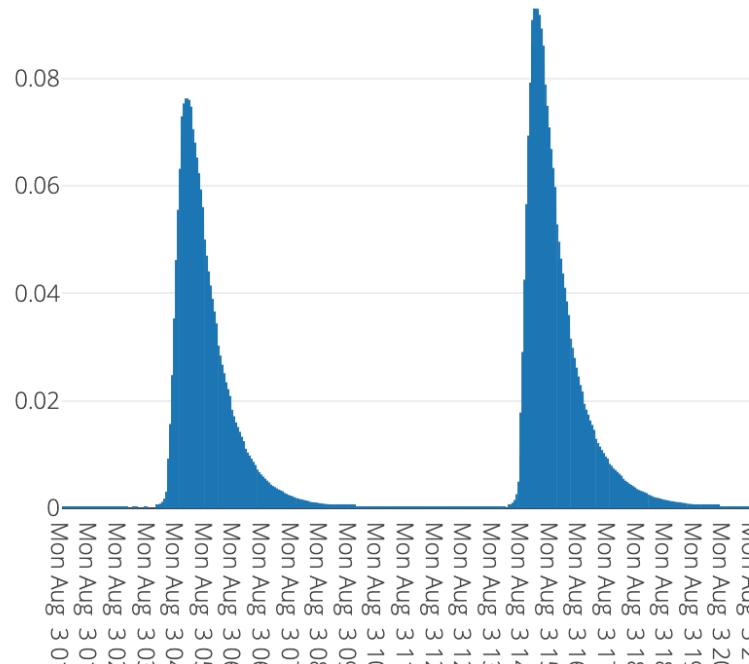
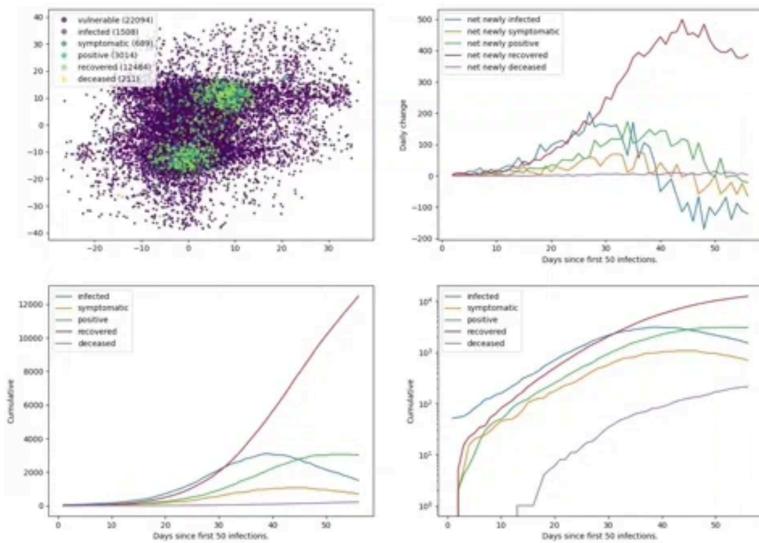
Five minute returns of bitcoin, ethereum, ripple, cardano and iota cryptocurrencies

Explanation: <https://www.linkedin.com/pulse/call-contributions-copula-contest-where-carefully-can-cotton-phd/>

What can you do with a microprediction API?

Surrogate models

https://www.microprediction.org/stream_dashboard.html?stream=pandemic_infected



Generative model: www.swarmprediction.com

How to drive a crawler

See MicroCrawler class

```
def include_sponsor(self, sponsor=None, **ignore):
    """ Override this as you see fit to select streams for your crawler """
    return True

def include_stream(self, name=None, **ignore):
    """ Override this as you see fit to select streams for your crawler
        For example your crawler might like z streams or not, or might require
        a minimum number of lags in the time series.
    """
    return True

    def candidate_streams(self):
        return ['three_body_x.json','three_body_y.json','three_body_z.json']

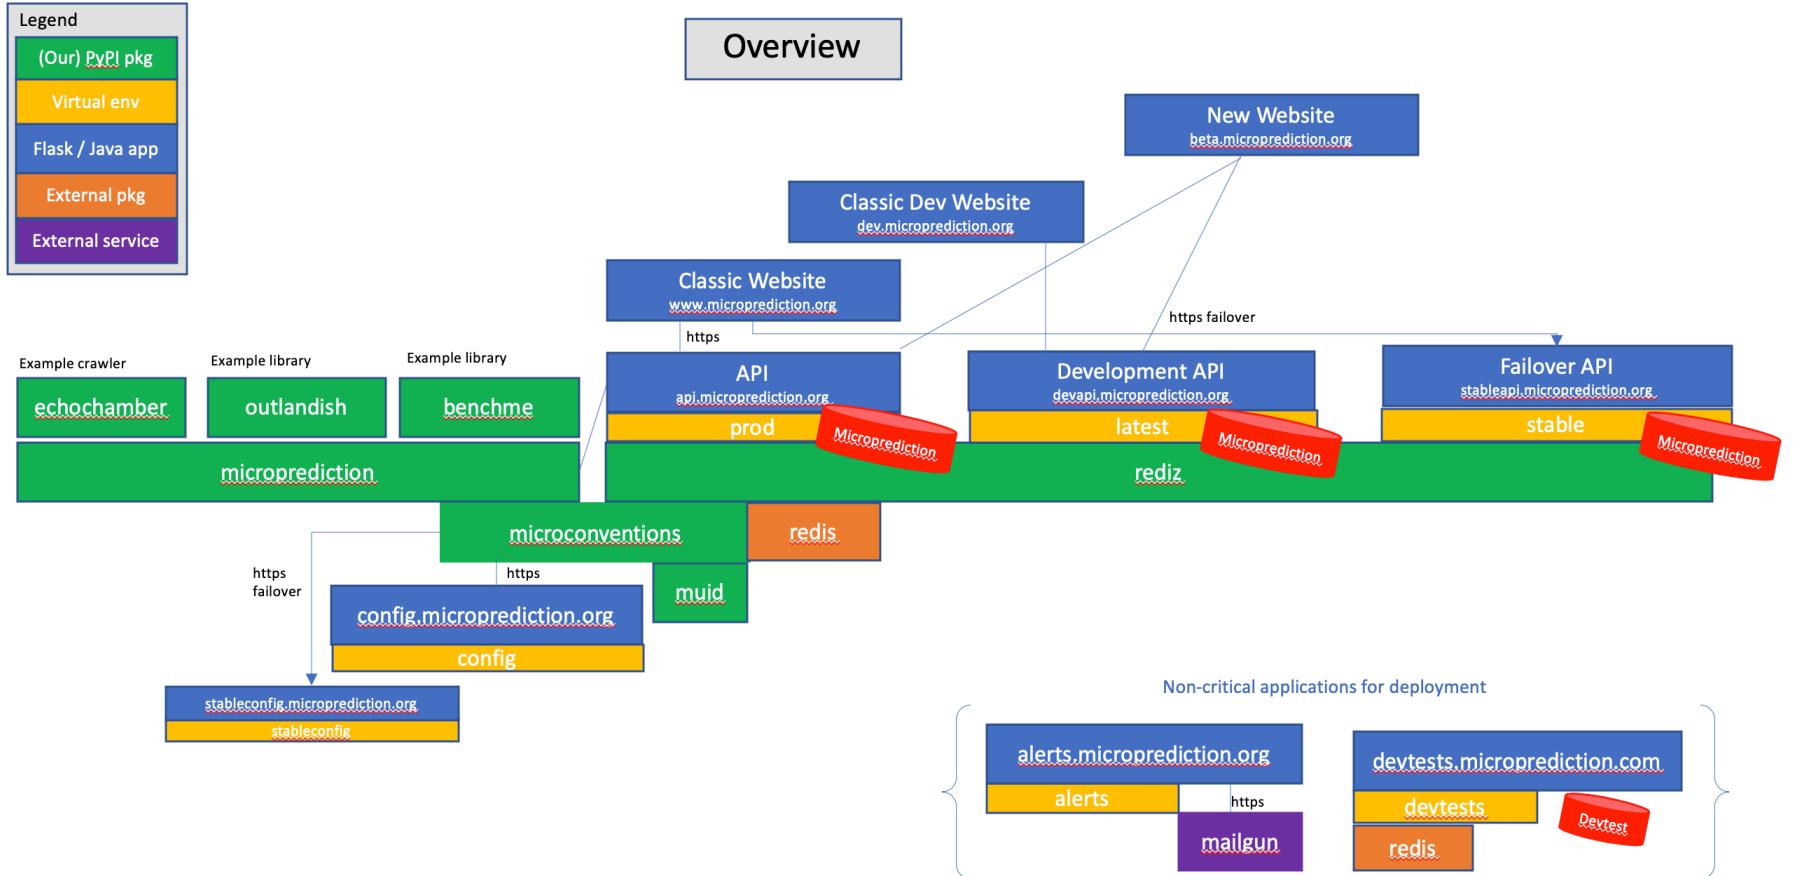
def include_delay(self, delay=None, name=None, **ignore):
    return True

    def downtime(self,seconds,**ignored):
        """ This will be called when a small gap opens up between times when you need to predict

```

<https://www.linkedin.com/pulse/economical-statistics-how-modify-prediction-crawlers-cotton-phd/>

Behind the scenes: architecture



HOW TO CONTRIBUTE

Dashboard

Leaderboard

Here are some suggestions if you would like to help.

1. Start [crawling](#) with thoughtful, well designed statistical algorithms.
2. Produce a live, clean JSON source of data of civic significance that changes at least five times an hour.
3. Write an R client.
4. Write a Julia client.
5. Make other suggestions by creating a [github issue](#).
6. Share or provide thoughtful feedback on [articles](#) on the topic to help get the word out.

<https://github.com/microprediction>



Eric Lou · 1st

CS + Math Student at Stanford University |

- Wrote the front end
- Winning crawlers



Rusty Conover · 1st

Experienced and innovative software executive with a track record of building businesses, platforms and applications.

- MUIDs in Java, Julia, Rust
- ZK-MUID proofs